

Photo Editor

Low-Level Design & Detailed Test Plan

Krishhiv Mehra & Siddhant Rau

10 November 2025

Document Purpose. This specification expands the previously submitted High-Level Design (Assignment 4) by detailing module internals (functions and classes), on-wire contracts, and a comprehensive test plan. Any architectural adjustments since, are folded directly into the relevant sections below.

1 Detailed Component & Module Design

1.1 Frontend - NextJS

- **Framework & State.** NextJS (App Router) with React Query for server state¹ and Zustand for editor UI state.
- **Key Routes.** `/library`, `/albums/[id]`, `/editor/[photoId]`, `/share/[token]` (public share with SSR²).
- **Core Components.** `Uploader`, `ImageCanvas` (Fabric.js/Konva adapter), `ToolSidebar`, `HistogramPanel`, `OpGraphInspector`, `ShareDialog`.
- **Editor State.** A pure JSON *operation graph* (“op-graph”) represents the sequence of edits. The canvas applies the same transforms client-side for instant preview; no pixels are persisted until *commit*.
- **Safe Retrieval.** Originals/variants are downloaded via signed Azure Blob SAS³ URLs. Public share pages are SSR with cache headers.
- **Validation.** Zod schemas mirror backend Pydantic models to prevent drift.

1.2 API Service (FastAPI)

Responsibilities. Auth, uploads, albums/photos, edit commit, search, sharing, and metadata (EXIF⁴/ICC⁵) handling.

Package Layout:

- `api/main.py` — app factory, middleware (CORS/HSTS), request ID logging.
- `api/deps.py` — auth dependencies, DB session management.

¹Server State: data fetched from remote APIs and cached on the client for synchronization.

²SSR: Server-Side Rendering.

³SAS: Shared Access Signature.

⁴EXIF: Exchangeable Image File Format.

⁵ICC: International Color Consortium (color profiles).

- `api/config.py` — typed settings via `pydantic-settings`.
- `api/routers/auth.py` — login/refresh (JWT⁶).
- `api/routers/uploads.py` — SAS issuance; chunked block-blob completion.
- `api/routers/albums.py` — CRUD + ACL⁷ checks.
- `api/routers/photos.py` — ingest, metadata read, listing, search.
- `api/routers/edits.py` — op-graph validation, commit, versioning.
- `api/routers/search.py` — tag/EXIF queries, pagination.
- `api/routers/shares.py` — share token issue/revoke, scope & expiry.
- `api/schemas/*.py` — Pydantic DTOs (`album.py`, `photo.py`, `version.py`, `share.py`, `errors.py`).
- `api/services/image_service.py` — Pillow/OpenCV transforms, tiled processing.
- `api/services/storage_service.py` — Azure Blob put/get via SAS.
- `api/services/version_service.py` — variant key derivation & metadata persistence.
- `api/services/exif_service.py` — EXIF read/strip GPS; ICC preservation.
- `api/repos/*.py` — SQLAlchemy ORM⁸ repositories for users, albums, photos, versions.
- `api/utils/opgraph.py` — schema, canonicalization, hashing.
- `api/utils/security.py` — JWT, share tokens (hashed at rest).
- `api/utils/idempotency.py` — Idempotency-Key enforcement for safe retries.

1.3 Database (PostgreSQL) – Schemas & Columns

We use PostgreSQL with SQLAlchemy 2.0. Naming: `snake_case` tables; `created_at/updated_at` timestamps; FKs with `ON DELETE CASCADE`. A partial GIN⁹ index supports tag search.

Schema: core

| Table | Column (type) | Why it exists |
|--------|--|--|
| users | id (uuid, PK) email (text, unique) password_hash (text) created_at (timestamp) tamptz) | Primary identity for auth & ownership. Unique login; lookup key. Store Argon2/BCrypt hash, not plaintext. Auditing and sort by recency. |
| albums | id (uuid, PK) owner_id (uuid, FK users.id) title (text) | Logical grouping of photos for users/teams. Ownership & ACL checks. Human-readable label. |

⁶JWT: JSON Web Token.

⁷ACL: Access Control List.

⁸ORM: Object-Relational Mapping.

⁹GIN: Generalized Inverted Index.

| Table | Column (type) | Why it exists |
|----------|---|--|
| | created_at (times-tz) | Ordering & lifecycle ops. |
| photos | id (uuid, PK) album_id (uuid, FK albums.id) blob_path (text) mime_type (text) width, height (int) exif (jsonb) created_at (times-tz) | Stable identity for an uploaded original. Belongs-to relation. Pointer to Blob URL path. Validate processing/serving. Prevent illegal ops; UI layout. Persist useful camera metadata. Traceability. |
| versions | id (uuid, PK) photo_id (uuid, FK photos.id) op_graph (jsonb) variant_key (text, unique) format (text) bytes (int) created_at (times-tz) | Each committed edit = a version. Back-reference to original. Audit the exact edits applied. Address immutable CDN ¹⁰ path. Serving/analytics; e.g., jpg/png/webp. Size tracking for quotas. Ordering/history. |
| shares | id (uuid, PK) photo_id (uuid, FK photos.id) scope (text) token_hash (text) expires_at (times-tz) | Public/limited access token record. What is being shared. view or comment etc. Store a hash, not raw token. Auto-expiry enforcement. |

Schema: tags

| Table | Column (type) | Why it exists |
|------------|---|--|
| photo_tags | photo_id (uuid, FK photos.id) tag (text) | Many-to-many mapping from photo to tag. Search & organization; GIN indexed. |

Indices & Constraints.

- users.email unique; versions.variant_key unique (*idempotency cache*).
- GIN index on photos.exif keys commonly queried; partial index on tags.photo_tags(tag).
- ON DELETE CASCADE for albums→photos and photos→versions.

¹⁰CDN: Content Delivery Network.

2 Operation Graph - Clear Explanation

2.1 What it is

An **operation graph** is a compact JSON description of *what* edits to apply and in *which order*. Think of it as a recipe; the original image plus this recipe renders the same output every time. This enables:

- **Reproducibility.** Same inputs ⇒ same result.
- **Caching.** We compute a short hash of the recipe to create a reusable *variant key*.
- **Auditability.** You can always see which edits produced a version.

2.2 Minimal Example

```
{  
  "photoId": "uuid-1234",  
  "ops": [  
    {"op": "crop", "x": 120, "y": 90, "w": 1024, "h": 768},  
    {"op": "rotate", "deg": 90},  
    {"op": "exposure", "ev": 0.33}  
,  
  "output": {"maxW": 2048, "format": "auto", "quality": "auto"}  
}
```

The engine applies crop then rotate then exposure, and finally chooses an output encoding (see Section 3).

2.3 Canonicalisation

To make caching reliable we canonicalise the JSON before hashing:

- 1) Round numeric fields to three decimals; normalize colors to #rrggb.
- 2) Freeze the operation order where it matters (e.g., exposure before contrast).
- 3) Remove irrelevant whitespace and sort object keys.

Result: identical edits *always* produce the same canonical JSON and therefore the same hash & filename.

2.4 Hash and Variant Key

We compute `hash = BLAKE3(photoId || canonical_json)` and set `variant_key = v/{photoId}/{hash[0..9]}.{ext}`. This makes every version addressable at an immutable path.

2.5 Supported Operations

- `crop`: $x, y \geq 0$; $w, h \geq 1$.
- `rotate`: degrees in $[-180, +180]$; bicubic resampling.
- `scale`: $\text{maxW}/\text{maxH} > 0$; Lanczos downscale; upscale capped at $2\times$.
- Tone/Color: `exposure`, `brightness`, `contrast`, `saturation`, `hue` within bounded amounts in linear RGB with ICC preserved.

- Blur/Sharpen: gaussian $\sigma \in [0, 5]$; unsharp mask (radius, percent, threshold).
- Text/Watermark: font allowlist; color contrast defaults meet AA¹¹.

3 Image Processing & Adaptive Encoding

- **Libraries.** Pillow/OpenCV; optional AVIF plugin. Tiled processing avoids excessive memory for very large images.
- **Color.** Adjust tone/contrast in linear RGB; convert back to sRGB; keep ICC; strip GPS from EXIF on public variants.
- **Formats.** JPEG (progressive, 4:2:0) for photos; WebP for size gains $\geq 15\%$; PNG or WebP-lossless for graphics/alpha.
- **Heuristic.** Probe on a 256px sample to select target format; encode full-res once.
- **Quality Gates.** Golden PSNR¹² ≥ 38 dB or SSIM¹³ ≥ 0.98 for auto-lossy; else fallback to lossless.

4 Custom Protocols & Contracts

- **Upload.** Client requests SAS; uploads via Azure Block Blob (4–8 MB blocks in parallel). Client posts a manifest to `/uploads/complete`; server verifies and persists metadata.
- **Edit Commit (Idempotent).** POST `/photos/{id}/edits` with op-graph JSON and `Idempotency-Key` header; server returns a new `version` with a signed URL.
- **Variant Retrieval.** Immutable GET to `v/{photoId}/{hash}.{ext}` with long-lived cache headers via CDN.

5 Low-Level Sequences

5.1 Commit Edit

- 1) Validate auth, ACL, and `photoId`; load photo metadata.
- 2) Validate & canonicalize op-graph; compute hash & variant key.
- 3) If variant already exists, persist only the `versions` row (idempotent fast-path).
- 4) Render pixels (tiled); choose encoding; upload variant to Blob.
- 5) Persist `versions` row and return DTO with signed URL.

5.2 Upload

- 1) Client obtains SAS for path `originals/{userId}/YYYY/MM/uuid.ext`.
- 2) Client uploads blocks with retries/backoff; completes with manifest.
- 3) API verifies, sniffs type/dimensions, writes `photos` row, pre-generates thumbnails.

¹¹WCAG AA: Web Content Accessibility Guidelines, Level AA.

¹²PSNR: Peak Signal-to-Noise Ratio.

¹³SSIM: Structural Similarity Index Measure.

6 Test Strategy & Quality Gates

6.1 Methodology

Test Pyramid: heavy unit & property tests; targeted integration; concise E2E¹⁴ via Playwright; separate performance and security suites. Environments: local (SQLite + Azurite), CI (PostgreSQL + Azurite), staging (Azure).

6.2 Unit Tests (PyTest + Hypothesis)

- **opgraph.** Canonicalisation (ordering, rounding), hashing determinism, invalid ranges rejected.
- **image_service.** Each op validated on synthetic images; rotation invariants; bounds.
- **version_service.** Stable key derivation; collision checks.
- **storage_service.** SAS scoping/expiry; MIME sniffing; type mismatch handling.
- **repos.** CRUD with constraint violations & rollbacks.

6.3 Integration Tests

- **Upload.** Chunked upload via Azurite; photos row exists; thumbnails reachable.
- **Edit Commit.** Given an original, post op-graph; confirm variant exists; bytes stable; ICC present; EXIF GPS stripped for public.
- **Search.** Tag + EXIF filters return expected sets with stable pagination.
- **Auth/ACL.** Share token scope respected; expiry enforced.

6.4 E2E

- Scenario: signup → upload → crop/rotate → save → open share link (mobile viewport).
- Accessibility: axe-core checks; keyboard navigation; dialog focus traps.

6.5 Performance & Reliability

- **Load.** Locust: 10→100 users editing distinct 24MP photos; record p50/p95; error rates.
- **Soak.** 1-hour continuous uploads of 10MB images with 1% network failure injection; verify no orphaned rows/objects.
- **Memory Ceiling.** Per-process RSS < 60% of VM during tiled operations.

6.6 Security Testing

- JWT expiry/forgery; refresh rotation; nonce checks.
- SAS abuse prevention (path scoping, minimal permissions); optional IP policy.
- OWASP vectors: malformed/oversized uploads; content-type confusion; SVG/script injection; CSP¹⁵ verified.
- Privacy: EXIF GPS stripping; no sensitive headers in responses.

¹⁴E2E: End-to-End.

¹⁵CSP: Content Security Policy.

7 Representative Test Cases

- **TC-U-001** Op-graph canonicalization \Rightarrow expect rounded JSON and stable hash.
- **TC-U-012** Rotate 90° on 8×8 checkerboard \Rightarrow pixel-perfect result.
- **TC-I-021** 50MB JPEG chunked upload (reordered blocks) \Rightarrow server reassembles; DB row created; thumbnails present.
- **TC-I-034** Edit commit idempotency with same `Idempotency-Key` \Rightarrow single version row; identical payload.
- **TC-E-005** Share link with past `expires_at` \Rightarrow 403 forbidden.
- **TC-P-009** p95 edit latency for 24MP crop+rotate+exposure $\leq 2.5\text{s}$ on B2s VM¹⁶.
- **TC-S-007** Malicious SVG upload \Rightarrow rejected with 415/422; sanitized preview only.

8 Requirement-to-Test Traceability

R-UP-01 (Chunked uploads) -> TC-I-021, E2E Scenario
R-ED-02 (Deterministic edits) -> TC-U-001, TC-I-034, golden-image suite
R-SH-03 (Share scopes & expiry) -> TC-E-005, E2E Scenario
R-PR-04 (Privacy: GPS stripping) -> Integration + Security tests
R-PF-05 (p95 $\leq 2.5\text{s}$) -> TC-P-009 (Locust run)

9 Operational Concerns

- **Observability.** Structured JSON logs (`reqId`, route, status, latency); optional Sentry; Azure Monitor counters.
- **Backups.** Nightly `pg_dump`; Blob lifecycle policy; quarterly restore drills.
- **Reconciliation.** Orphan sweeper for DB \leftrightarrow Blob mismatches with dry-run reports.

ChatGPT 5 was used to research and for reference purposes.

¹⁶VM: Virtual Machine.